

▼ Common Test I. Multi-Class Classification

Task: Build a model for classifying the images into lenses using PyTorch or Keras. Pick the most appropriate approach and discuss your strategy.

Dataset: [dataset.zip - Google Drive](#)

Dataset Description: The Dataset consists of three classes, strong lensing images with no substructure, subhalo substructure, and vortex substructure. The images have been normalized using min-max normalization, but you are free to use any normalization or data augmentation methods to improve your results.

Evaluation Metrics: ROC curve (Receiver Operating Characteristic curve) and AUC score (Area Under the ROC Curve)

▼ Importing required libraries

```
1 #Utilities
2 import os
3 import gc
4 import glob
5 import numpy as np
6 import pandas as pd
7 from tqdm.notebook import tqdm
8
9 #Loading image and plotting visualizations/images
10 from PIL import Image
11 import seaborn as sns
12 import matplotlib.pyplot as plt
13
14 #PyTorch framework
```

```

15 import torch
16 import torch.nn as nn
17 import torch.optim as optim
18 from torch.optim.lr_scheduler import CosineAnnealingWarmRestarts
19 from torch.utils.data import DataLoader, Dataset
20 from torchvision import utils
21
22 #Evaluation metrics
23 from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve, auc
24
25 #For pre-trained model
26 import sys
27 sys.path.append('../input/timm-pytorch-image-models/pytorch-image-models-master')
28 import timm
29
30 np.random.seed(7)
31 torch.manual_seed(7)
32
33 device='cuda' if torch.cuda.is_available() else 'cpu'
34 device

```

'cuda'

▼ Creating a custom dataset class

```

1 class CustomDataset(Dataset):
2     def __init__(self, root_dir, transform = None):
3         root_list = glob.glob(root_dir)
4         self.class_map = {}
5         self.class_distribution = {}
6         self.transform = transform
7
8         for img_path in root_list:
9             class_name = img_path.split(os.sep)[-2]
10            if class_name not in self.class_distribution:
11                self.class_map[class_name] = []

```

```
11         self.class_distribution[class_name] = 1
12     else:
13         self.class_distribution[class_name] +=1
14
15     for index, entity in enumerate(self.class_distribution):
16         self.class_map[entity] = index
17     print("Dataset Distribution:\n")
18     print(self.class_distribution)
19     print("\n\nClass indices:\n")
20     print(self.class_map)
21
22     self.data = []
23     for img_path in tqdm(root_list):
24         class_name = img_path.split(os.sep)[-2]
25         self.data.append([img_path, class_name])
26
27     def __len__(self):
28         return len(self.data)
29
30     def __getitem__(self, idx):
31         img_path, class_name = self.data[idx]
32         img = np.load(img_path)
33         img = torch.tensor(img, dtype=torch.float)
34
35         class_id = self.class_map[class_name]
36         class_id = torch.tensor(class_id)
37
38         return img, class_id


1 #Using a batch size of 128
2 BS = 128


1 train_path = r'../input/ml4sci-deeplense-commontask/dataset/train/*/*'
2 train_dataset = CustomDataset(train_path)
3
4 val_path = r'../input/ml4sci-deeplense-commontask/dataset/val/*/*'
```

```

5 val_dataset = CustomDataset(val_path)
6
7 print(len(train_dataset), len(val_dataset))

```

Dataset Distribution:

```
{'no': 10000, 'vort': 10000, 'sphere': 10000}
```

Class indices:

```
{'no': 0, 'vort': 1, 'sphere': 2}
0%|          | 0/30000 [00:00<?, ?it/s]
```

Dataset Distribution:

```
{'no': 2500, 'vort': 2500, 'sphere': 2500}
```

Class indices:

```
{'no': 0, 'vort': 1, 'sphere': 2}
0%|          | 0/7500 [00:00<?, ?it/s]
30000 7500
```

▼ Creating data loaders separately for train data and val/test data

```

1 train_loader = DataLoader(train_dataset, batch_size = BS, shuffle = True)
2 val_loader = DataLoader(val_dataset, batch_size = BS, shuffle = False)

```

```


1 single_batch = next(iter(train_loader))
2 print(f"The dimensions of a single batch is {single_batch[0].shape}")


```


```
The dimensions of a single batch is torch.Size([128, 1, 150, 150])
```

▼ Plotting a shuffled single batch from the train loader

```
1 single_batch_grid = utils.make_grid(single_batch[0], nrow=8)
2 plt.figure(figsize = (20,70))
3 plt.imshow(single_batch_grid.permute(1, 2, 0))
4 plt.savefig("Single_batch_of_train_loader.png")
```







▼ Creating model

Using the Inception Resnet V2 pre-trained model as the backbone, adding my own classifier head and fine-tuning to this dataset

```
1 class pre_trained_model(nn.Module):
2
3     def __init__(self, pretrained = True):
4         super().__init__()
5         self.model = timm.create_model('inception_resnet_v2', pretrained = pretrained, in_chans = 1)
6         # num_in_features = self.model.get_classifier().in_features
7
8         for param in self.model.parameters():
9             param.requires_grad = True
10
11         self.fc = nn.Sequential(
12             nn.Linear(1536 * 3 * 3, 1024),
13             nn.PReLU(),
14             nn.BatchNorm1d(1024),
15             nn.Dropout(p = 0.5),
```



```

16
17         nn.Linear(1024, 128),
18         nn.PReLU(),
19         nn.BatchNorm1d(128),
20         nn.Dropout(p = 0.5),
21
22         nn.Linear(128, 3)
23     )
24
25     def forward(self, x):
26         x = self.model.forward_features(x)
27         x = x.view(-1, 1536 * 3 * 3)
28         x = self.fc(x)
29         return x

```

```

1 model = pre_trained_model()
2
3 #Verifying output of model
4
5 x = torch.randn(128, 1, 150, 150)
6 model(x).shape

```

Downloading: "https://github.com/rwightman/pytorch-image-models/releases/download/v0.1-weights/inception_resnet_v2-940b1cd6.pth"
 torch.Size([128, 3])



```

1 def calculate_accuracy(y_pred, y_truth):
2     y_pred_softmax = torch.log_softmax(y_pred, dim = 1)
3     _, y_pred_labels = torch.max(y_pred_softmax, dim = 1)
4
5     correct_preds = (y_pred_labels == y_truth).float()
6     acc = correct_preds.sum() / len(correct_preds)
7     acc = torch.round(acc*100)
8
9     return acc

```

```
1 def train_epoch(model, dataloader, criterion, optimizer):
2     model.train()
3     train_loss = []
4     train_accuracy = []
5
6     loop=tqdm(enumerate(dataloader),total = len(dataloader))
7
8     for batch_idx, (img_batch,labels) in loop:
9
10         X = img_batch.to(device)
11         y_truth = labels.to(device)
12
13         #forward prop
14         y_pred = model(X)
15
16         #loss and accuracy calculation
17         loss = criterion(y_pred, y_truth)
18         accuracy = calculate_accuracy(y_pred, y_truth)
19
20         #backprop
21         optimizer.zero_grad()
22         loss.backward()
23         optimizer.step()
24
25         #batch loss and accuracy
26         # print(f'Partial train loss: {loss.data}')
27         train_loss.append(loss.detach().cpu().numpy())
28         train_accuracy.append(accuracy.detach().cpu().numpy())
29
30     return model, np.mean(train_loss), np.mean(train_accuracy)
```

```
1 def val_epoch(model, dataloader,criterion):
2     model.eval()
3     val_loss = []
4     val_accuracy = []
5
```

```
6     with torch.no_grad():
7
8         loop=tqdm(enumerate(dataloader),total=len(dataloader))
9
10        for batch_idx, (img_batch,labels) in loop:
11            X = img_batch.to(device)
12            y_truth = labels.to(device)
13
14            #forward prop
15            y_pred = model(X)
16
17            #loss and accuracy calculation
18            loss = criterion(y_pred, y_truth)
19            accuracy = calculate_accuracy(y_pred, y_truth)
20
21
22            #batch loss and accuracy
23            # print(f'Partial train loss: {loss.data}')
24            val_loss.append(loss.detach().cpu().numpy())
25            val_accuracy.append(accuracy.detach().cpu().numpy())
26
27        return np.mean(val_loss), np.mean(val_accuracy)


1 def fit_model(model,criterion,optimizer):
2     loss_dict = {'train_loss':[],'val_loss':[]}
3     acc_dict = {'train_accuracy':[],'val_accuracy':[]}
4
5     for epoch in range(EPOCHS):
6         print(f"Epoch {epoch+1}/{EPOCHS}:")
7         model, train_loss, train_accuracy = train_epoch(model, train_loader, criterion, optimizer)
8         val_loss, val_accuracy = val_epoch(model, val_loader, criterion)
9
10        print(f'Train loss:{train_loss}, Val loss:{val_loss}')
11        loss_dict['train_loss'].append(train_loss)
12        loss_dict['val_loss'].append(val_loss)
13        print(f'Train accuracy: {train_accuracy}, Val accuracy:{val_accuracy}')
14        acc_dict['train_accuracy'].append(train_accuracy)
```

```
15         acc_dict['val_accuracy'].append(val_accuracy)
16
17
18     return model, loss_dict, acc_dict
```

▼ Initializing the model and deciding the hyperparameter values

Multi-class classification problem -> Loss function : CrossEntropyLoss

Model trained for 10 epochs

Adam Optimizer used with learning rate $3e-4$

```
1 model = pre_trained_model().to(device)
2
3 criterion = nn.CrossEntropyLoss()
4 EPOCHS = 10
5 LR = 3e-4
6
7 optimizer = optim.Adam(model.parameters(),lr=LR)

1 #Training model
2 model, loss_dict, acc_dict = fit_model(model,criterion,optimizer)
```

```
Epoch 1/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.9953879714012146, Val loss:0.9294667840003967
Train accuracy: 51.11489486694336, Val accuracy:58.779659271240234
Epoch 2/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.5411764979362488, Val loss:0.45875853300094604
Train accuracy: 77.87659454345703, Val accuracy:81.66101837158203
Epoch 3/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.35474199056625366, Val loss:0.3722762167453766
Train accuracy: 86.22553253173828, Val accuracy:85.30508422851562
Epoch 4/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.2617902457714081, Val loss:0.36338678002357483
Train accuracy: 90.2170181274414, Val accuracy:87.40677642822266
Epoch 5/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.19792300462722778, Val loss:0.4865666925907135
Train accuracy: 92.77021026611328, Val accuracy:80.16949462890625
Epoch 6/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.1542099267244339, Val loss:0.3650676906108856
Train accuracy: 94.31063842773438, Val accuracy:88.5762710571289
Epoch 7/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.11358343064785004, Val loss:0.38710132241249084
Train accuracy: 95.94893646240234, Val accuracy:86.81356048583984
Epoch 8/10:
 0%|          | 0/235 [00:00<?, ?it/s]
 0%|          | 0/59 [00:00<?, ?it/s]
Train loss:0.09248575568199158, Val loss:0.39652055501937866
```

```
1 #saving the trained model
```

```

2 PATH = "inception_resnetV2_finetuned.pth"
3 torch.save(model.state_dict(), PATH)

# Save the state dict to a file
# .....

1 #deleting the trained model instance, creating a new one and loading the saved train model
2 del model
3 gc.collect()
4
5 model = pre_trained_model().to(device)
6 model.load_state_dict(torch.load(PATH))

<All keys matched successfully>


```

▼ Loss through the epochs

```

1 # Plot losses
2 plt.figure(figsize=(9,7))
3 plt.semilogy(loss_dict['train_loss'], label='Train')
4 plt.semilogy(loss_dict['val_loss'], label='Valid')
5 plt.xlabel('Epoch')
6 plt.ylabel('Average Loss')
7 #plt.grid()
8 plt.legend()
9 #plt.title('loss')
10 plt.show()
11 plt.savefig("Loss_history.png")

```




▼ Val/test accuracy through the epochs

```

1 # Plot accuracy
2 plt.figure(figsize=(9,7))
3 plt.semilogy(acc_dict['train_accuracy'], label='Train')
4 plt.semilogy(acc_dict['val_accuracy'], label='Valid')
5 plt.xlabel('Epoch')
6 plt.ylabel('Average Accuracy')
7 #plt.grid()
8 plt.legend()
9 #plt.title('loss')
10 plt.show()
11 plt.savefig("Accuracy_history.png")

```



▼ Final prediction

(Usually this is done on another separate test set. In this case, I have considered the test and val sets to be same)

Figure size 432x288 with 0 axes

```

1 def test_epoch(model, dataloader, criterion):
2
3     model.eval()
4     test_loss = []
5     test_accuracy = []
6
7     y_pred_list = []
8     y_truth_list = []
9     y_pred_prob_list = []
10
11     with torch.no_grad():
12
13         loop=tqdm(enumerate(dataloader),total=len(dataloader))

```



```

14
15     for batch_idx, (img_batch, labels) in loop:
16         X = img_batch.to(device)
17         y_truth = labels.to(device)
18         y_truth_list.append(y_truth.detach().cpu().numpy())
19
20         #forward prop
21         y_pred = model(X)
22         y_pred_softmax = torch.log_softmax(y_pred, dim = 1)
23         y_pred_prob_list.append(y_pred_softmax.detach().cpu().numpy())
24         _, y_pred_labels = torch.max(y_pred_softmax, dim = 1)
25         y_pred_list.append(y_pred_labels.detach().cpu().numpy())
26
27         #loss and accuracy calculation
28         loss = criterion(y_pred, y_truth)
29         accuracy = calculate_accuracy(y_pred, y_truth)
30
31
32         #batch loss and accuracy
33         # print(f'Partial train loss: {loss.data}')
34         test_loss.append(loss.detach().cpu().numpy())
35         test_accuracy.append(accuracy.detach().cpu().numpy())
36
37     return y_pred_prob_list, y_pred_list, y_truth_list, np.mean(test_loss), np.mean(test_accuracy)

1 y_pred_prob_list, y_pred_list, y_truth_list, test_loss, test_accuracy = test_epoch(model, val_loader, criterion)
2
3 print(test_loss, test_accuracy)

0%|          | 0/59 [00:00<?, ?it/s]
0.49393848 86.932205

```

An accuracy of 89.86% has been achieved on the test set

1 #To flatten the outputs since the predictions are generated in batches w.r.t the data loader

```

2
3 def flatten_list(x):
4     flattened_list = []
5     for i in x:
6         for j in i:
7             flattened_list.append(j)
8
9     return flattened_list

1 y_pred_list_flattened = flatten_list(y_pred_list)
2 y_truth_list_flattened = flatten_list(y_truth_list)
3 y_pred_prob_list_flattened = flatten_list(y_pred_prob_list)

1 idx2class = {v: k for k, v in train_dataset.class_map.items()}
2 class_names = [i for i in train_dataset.class_map.keys()]
3 idx2class

{0: 'no', 1: 'vort', 2: 'sphere'}

```

▼ Classification Report on the test set

```
1 print(classification_report(y_truth_list_flattened, y_pred_list_flattened, target_names = class_names))
```


	precision	recall	f1-score	support
no	0.89	0.87	0.88	2500
vort	0.92	0.85	0.88	2500
sphere	0.81	0.89	0.85	2500
accuracy			0.87	7500
macro avg	0.87	0.87	0.87	7500
weighted avg	0.87	0.87	0.87	7500

```
1 print(confusion_matrix(y_pred_list_flattened, y_truth_list_flattened))
```

```
[[2179   92  172]
 [   77 2117  102]
 [   244  291 2226]]
```

▼ Confusion Matrix from test set predictions

```
1 confusion_matrix_df = pd.DataFrame(confusion_matrix(y_truth_list_flattened, y_pred_list_flattened)).rename(columns=idx2class, index=idx2class)
2 fig, ax = plt.subplots(figsize=(19,12))
3 sns.heatmap(confusion_matrix_df, fmt = ".0f", annot=True, ax=ax)
4 plt.savefig("Confusion_matrix.png")
```



▼ Plotting ROC curve

```

1 #One-hot-encoding ground truths
2 temp_test_y = []
3
4 for i in range(len(y_truth_list_flattened)):
5     a = [0, 0, 0]
6     a[y_truth_list_flattened[i]] = 1
7     temp_test_y.append(a)
8
9 temp_test_y = np.array(temp_test_y)
10 temp_test_y.shape

```

```
(7500, 3)
```

```

1 #Note: The test set was not shuffled
2 temp_test_y[0:5]

```

```

array([[1, 0, 0],
       [1, 0, 0],
       [1, 0, 0],

```


```
[1, 0, 0],  
[1, 0, 0]])
```

```
1 y_pred_prob_list_flattened = np.array(y_pred_prob_list_flattened)  
2 y_pred_prob_list_flattened.shape
```

```
(7500, 3)
```

```
1 fpr = dict()  
2 tpr = dict()  
3 roc_auc = dict()  
4  
5 for i in range(3):  
6     fpr[i], tpr[i], _ = roc_curve(temp_test_y[:, i], y_pred_prob_list_flattened[:, i])  
7     roc_auc[i] = auc(fpr[i], tpr[i])
```

```
1 colors = ['red', 'blue', 'green']  
2 plt.figure(figsize = (19, 12))  
3  
4 for i, color in zip(range(3), colors):  
5     plt.plot(fpr[i], tpr[i], color=color, lw=2, label='ROC curve of class {0} (area = {1:0.2f})' ''.format(i, roc_auc[i]))  
6  
7 plt.plot([0, 1], [0, 1], 'k--', lw=2)  
8 plt.xlim([0.0, 1.0])  
9 plt.ylim([0.0, 1.05])  
10 plt.xlabel('False Positive Rate')  
11 plt.ylabel('True Positive Rate')  
12 plt.title('Receiver Operating Characteristic (Inception_Resnet_V2 model fine-tuned)')  
13 plt.legend(loc="lower right")  
14 plt.show()  
15 plt.savefig("ROC_curve.png")
```



<Figure size 432x288 with 0 Axes>

▼ ROC-AUC score